




Threading


Reese Hatfield



0

Threading

- Remember the Event Dispatch Thread?
- Let your Swing code run in parallel





Threading

- Java lets you write arbitrary code that runs on a separate thread(s)
- Two primary ways of doing this
 - Thread Class
 - Runnable interface






Threading

- Anytime you want to write some code that will run asynchronously
 - Must at least *start* in a public void run() method



Threading


- Override the run() method from the Thread class
- Requires you to extend Thread





Threading

- You implement the runnable interface
- Pass that into a Thread instance





Threading

- Let's try doing these
 - Gonna focus on extending from Thread
 - Likely what you'll need for Project 4





Threading

- Threads are created like any other Object
 - Thread t1 = new CustomThread
- Do *NOT* call the run() method directly
 - Start threads via their .start() method
- Do *NOT* kill a thread directly
 - Wait for them to finish





Threading

- When you start() a thread:
 - The JVM will handle calling the run method in the background
 - Your code will continue to run after
 - It will *not* wait until it is finished





Threading

- When you join() a thread:
 - Your code will completely stop
 - Until that thread has finished all of its work
 - (block until run is finished)





Threading

- With just being able to
 - Create a thread
 - Start a thread
 - Join a thread
- You can do *basically* anything* in parallel





Threading


- How would threads work with the *static* keyword?
- Static = bound to class, not instance of class
- Effectively a single instance





Threading

- CPUs are *really* effective at quick data lookups
 - This is one of their main purposes
- They will often cache data into registers





Threading

- Each thread can be thought of as a virtual CPU
- But if we want to share data across threads
 - A cache would cause our data be bad
 - Even if it's static
- To fix this, we use the *volatile* keywords
 - Often alongside the *static* keyword





Threading

- We saw how to *use* volatile
 - And when to use it
 - But what happens if we **don't**
- Lets try it:
 - Spawn a thread → do some work for a time
- After a second
 - Stop it from doing work in the main thread
 - Control work from a boolean





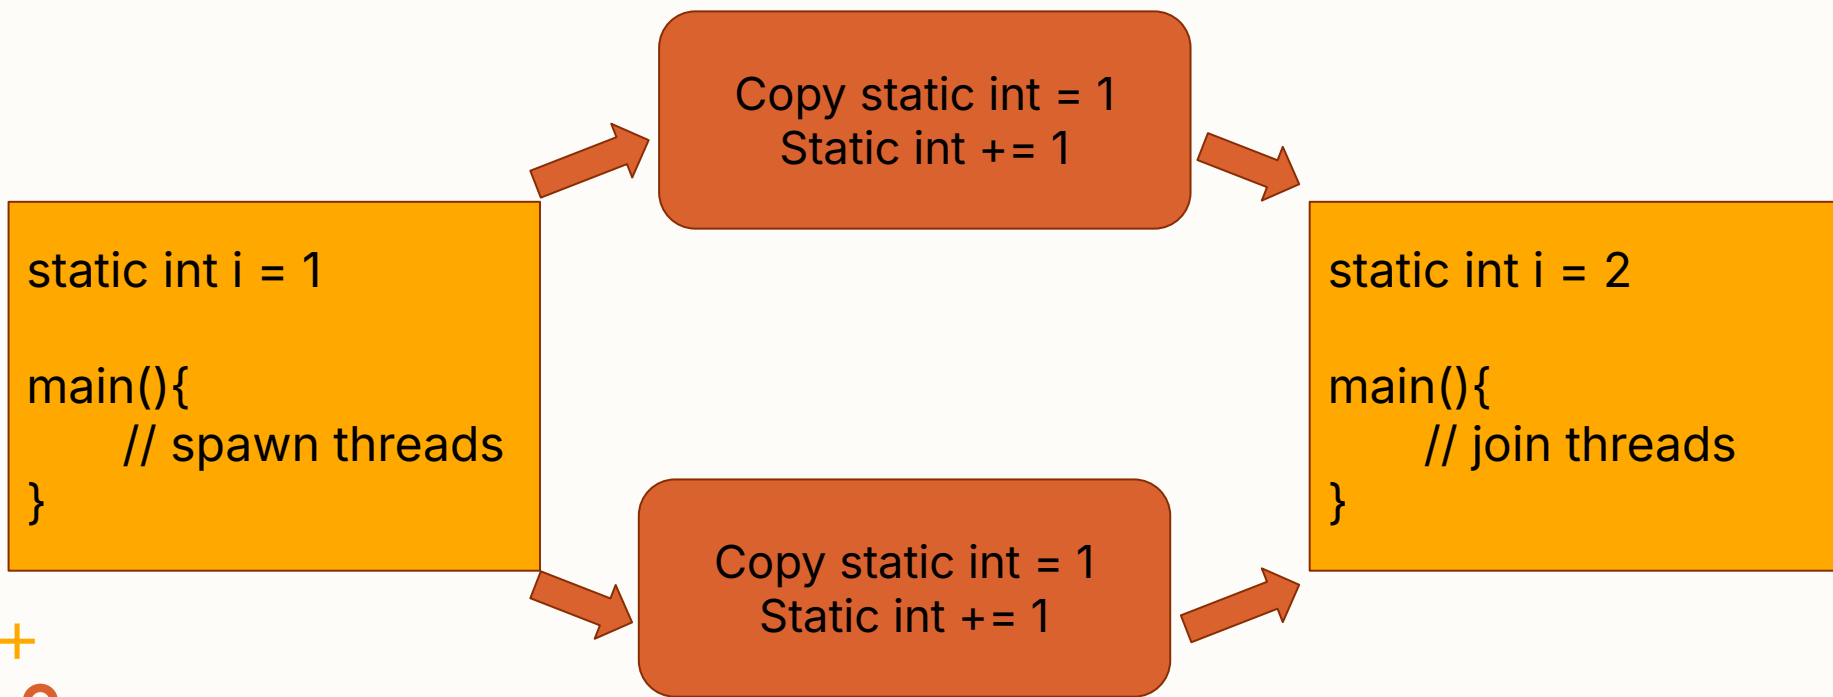
Threading

- Why does this break without *volatile*?
- A Thread gets its own environment
 - Separate variables, callstack, etc
- "Thread" will cache the value of "running"
- We need to tell it *not* to cache that value





Threading





Threading

- When should I use threading?
 - I/O bound operations
 - Files, etc
 - CPU bound operations
 - Must be "independent" operations





I/O Bound Threading


- I/O bound operations
- Reading/Writing to Files
 - Incredibly slow → see pyramid
 - Make multiple files
 - Have different threads access copies





I/O Bound Threading


- Files can only handle one operation at a time
- Threads must wait on each other





I/O Bound Threading

- To fix this bottleneck:
 - Copy the files for each thread





CPU Bound Threading

- CPU are really good at executing sequential instructions
- Sequence = Ordered
- But we need thing to run in parallel*





CPU Bound Threading

- Most things cannot inherently be parallelized
- Most, not all
- How do we know if a task can be parallelized?





CPU Bound Threading

- Task can be done concurrently if:
- Task is sufficiently independent
- One piece does not depend on the rest
 - Taking a sum $\Rightarrow A + B = B + A$
 - Searching in a list





CPU Bound Threading

- Threading is also often used for long, blocking tasks
- If you were doing an operation, just once, but that operation took forever...
- Throw that task on another thread





Thread Concerns

- We saw before that
 - Need* a volatile variable to share data
 - Usually statically access
- This ensures each thread uses the same data
 - Why did we need this?





Thread Concerns

- Each thread would cache the variables value without volatile
- Is there ever a better way to do this?
- What if we wanted to sync our threads at the method level
 - Instead of variable access level





Thread Concerns

- There is a keyword for this
 - synchronized
- Prevent threads from interfering with inconsistent memory
- Let's see this in action
 - Counter Class
 - Synchronized vs unsynchronized methods





Thread Safety

- By making our accessors synchronized
 - We made this class "thread safe"
 - Meaning, that our class state is always safe to access, even among different threads
- You'll frequently see this written in documentation





Thread Safety

- Thread safety ensure our memory is safe to access
- If different threads did not promise this, whoever reached the data last, will have the wrong result





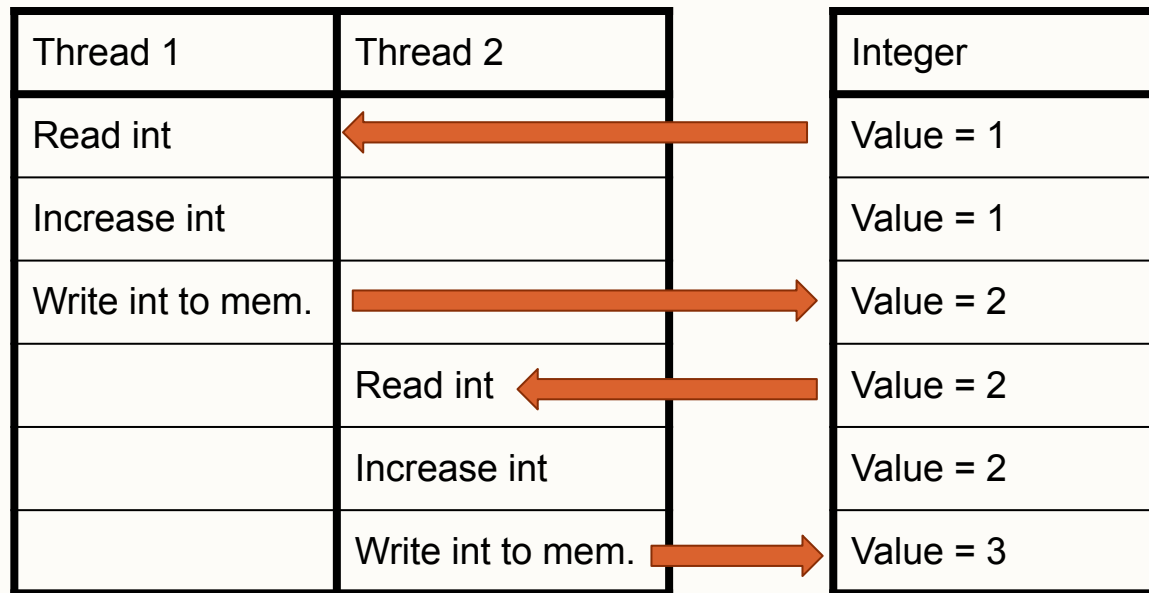
Thread Safety

- This is called a Race Condition
- Let's visualize this a little better
- CPU's execute sequential instructions
 - Read a value
 - Modify a value
 - Write to a value






Thread Safety





Thread Safety





Thread Safety

- The threads here ran at the same type
 - By sort of jumbling the timing of their instructions
 - This is a slightly more accurate depiction of what goes on
- Prevented by synchronized + volatile





Patterns

- The most common pattern you'll see in threading is "chunking"
 - Divide data into *variable* size chunks
 - Control for the number of threads
- Let's try it!

